

```
In [ ]: import pandas as pd
from itertools import combinations
from statsmodels.tsa.stattools import coint, adfuller
import matplotlib.pyplot as plt
import seaborn as sns; sns.set(style="whitegrid")
import statsmodels.api as sm
import numpy as np
import os
from xgboost import XGBClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import LabelEncoder
import warnings
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import RandomizedSearchCV
from xgboost import XGBRegressor
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from sklearn.model_selection import GridSearchCV
import xgboost as xgb
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd
from tabulate import tabulate
```

```
# Suppress SettingWithCopyWarning
warnings.filterwarnings("ignore", category=pd.errors.SettingWithCopyWarning)

# Suppress FutureWarning for deprecated functions
warnings.filterwarnings("ignore", category=FutureWarning)
```

```
In [ ]: sp500_tickers_df = pd.read_html('https://en.wikipedia.org/wiki/List_of_S&P_500_tickers')
sp500_tickers = sp500_tickers_df['Symbol'].tolist()
```

```
In [ ]: brim_tickers = ['ADBE', 'AMZN', 'AWR', 'BEN', 'CF', 'CNX', 'COG', 'CRM', 'DRI', 'F', 'FB', 'GOOGL', 'HPE', 'IBM', 'INTC', 'JPM', 'KROGER', 'L', 'MCD', 'MSFT', 'NKE', 'NFLX', 'ORCL', 'PFE', 'PG', 'PYPL', 'QCOM', 'R', 'ROST', 'T', 'TSLA', 'V', 'VZ', 'W', 'WAT', 'WMT', 'XOM']
brim_names = ['Adobe', 'Amazon', 'American States Water', 'Franklin Resources', 'Hewlett-Packard', 'IBM', 'Intel', 'JPMorgan Chase', 'Kroger', 'McDonald's', 'Microsoft', 'Nike', 'Netflix', 'Oracle', 'Pfizer', 'Procter & Gamble', 'Pyramid Energy Services', 'Ralph Lauren', 'Rite Aid', 'Ross Stores', 'Salesforce', 'Sealed Air', 'Shutterstock', 'Spotify', 'Target', 'T-Mobile', 'United Therapeutics', 'Verizon', 'Veeva Systems', 'Walmart', 'Wendy's', 'Xerox']

print(len(brim_names), len(brim_tickers))
```

30 30

```
In [ ]: current_directory = os.getcwd()
print("Current Directory:", current_directory)
datasets = {}
```

```

test_datasets = {}
tickers = []
for t in (sp500_tickers + brim_tickers):
    #path = "/Users/rashimohta/Downloads/sp500_data/" + t + ".csv"
    path = current_directory+"/sp500_data/" + t + ".txt" #use .csv or .txt
    #print("PATH: ",path)
    df = pd.read_csv(path)
    df["DT Date"] = pd.to_datetime(df['Date']) #add a column with the date

    if 'DT Date' in df.index.names and not 'DT Date' in df.columns:
        df['DT Date'] = df.index.copy()
        print("Added Column DT Date from Index")

    if len(df)==0:
        print("Data not available for ticker: ", t)
    #elif len(df) != 2516:
    elif len(df) < 100:
        print(f"{t} Removed from the stock universe : Fewer than 100 days data")
    else:
        datasets[t] = df
        test_datasets[t] = df
        tickers.append(t)
        """datasets[t] = df[:600]
        test_datasets[t] = df[-600:]
        tickers.append(t)"""

```

Current Directory: /Users/rashimohta/Downloads/Pairs Trading Project

Data not available for ticker: AMTM

Data not available for ticker: BRK.B

Data not available for ticker: BF.B

GEV Removed from the stock universe : Fewer than 100 days data

Data not available for ticker: SW

SOLV Removed from the stock universe : Fewer than 100 days data

```

In [ ]: current_directory = os.getcwd()
print("Current Directory:", current_directory)
datasets = {}
test_datasets = {}
tickers = []
for t in (sp500_tickers + brim_tickers):
    #path = "/Users/rashimohta/Downloads/sp500_data/" + t + ".csv"
    path = current_directory+"/sp500_data/" + t + ".txt" #use .csv or .txt
    #print("PATH: ",path)
    df = pd.read_csv(path)
    df["DT Date"] = pd.to_datetime(df['Date']) #add a column with the date

    if 'DT Date' in df.index.names and not 'DT Date' in df.columns:
        df['DT Date'] = df.index.copy()
        print("Added Column DT Date from Index")

    if len(df)==0:
        print("Data not available for ticker: ", t)
    #elif len(df) != 2516:
    elif len(df) < 100:

```

```

print(f"{t} Removed from the stock universe : Fewer than 100 days
else:
    datasets[t] = df
    test_datasets[t] = df
    tickers.append(t)
    """datasets[t] = df[:600]
    test_datasets[t] = df[-600:]
    tickers.append(t)"""

```

Current Directory: /Users/rashimohta/Downloads/Pairs Trading Project

Data not available for ticker: AMTM

Data not available for ticker: BRK.B

Data not available for ticker: BF.B

GEV Removed from the stock universe : Fewer than 100 days data

Data not available for ticker: SW

SOLV Removed from the stock universe : Fewer than 100 days data

```

In [ ]: brim_stock_pairs = [['BEN', 'COG'], ['DISCA', 'RIG'], ['DISCK', 'RIG'], [
    ['CNX', 'HBI'], ['AMZN', 'CRM'], ['MA', 'VFC'], ['FCX', 'G
    ['FCX', 'HBI'], ['DISCK', 'ESV'], ['DISCK', 'NE'], ['DISCA
    ['NBL', 'RIG'], ['CNX', 'GNW'], ['COG', 'DO'], ['HBI', 'NB
    ['DISCA', 'MA'], ['DISCK', 'MA'], ['RIG', 'RRC'], ['CF', '
    ['NE', 'RRC'], ['ADBE', 'RHT'], ['MA', 'RIG'], ['NBL', 'SW
    ['AWR', 'WTR'], ['SLB', 'PFE']]

```

```

In [ ]: def calculate_features(df):
    features = pd.DataFrame(index=df.index)

    features['Current Spread'] = df['Spread']
    features['Spread Returns'] = features['Current Spread'].pct_change()

    for days in [15, 10, 7, 5]:
        features[f'Sp Mean {days}days'] = features['Current Spread'].roll
        features[f'Sp/SpMean {days}days'] = features['Current Spread'] /

    column_order = ['Current Spread', 'Spread Returns'] + \
        [f'Sp Mean {days}days' for days in [15, 10, 7, 5]] + \
        [f'Sp/SpMean {days}days' for days in [15, 10, 7, 5]]

    features = features[column_order]

    return features

def prepare_data(datasets, pair):
    stock1, stock2 = pair

    df1 = datasets[stock1]
    df2 = datasets[stock2]

    if not isinstance(df1.index, pd.DatetimeIndex):
        df1['DT Date'] = pd.to_datetime(df1['DT Date'])
        df1.set_index('DT Date', inplace=True)

    if not isinstance(df2.index, pd.DatetimeIndex):

```

```

df2['DT Date'] = pd.to_datetime(df2['DT Date'])
df2.set_index('DT Date', inplace=True)

# Calculate spread
spread_df = pd.DataFrame(index=df1.index)
spread_df['Spread'] = df1['Adj Close'] - df2['Adj Close']

# Calculate features
features = calculate_features(spread_df)

# Calculate target variable (next day's spread return)
features['Target'] = features['Spread Returns'].shift(-1)
features['Pair'] = f"{pair[0]}_{pair[1]}"
return features.dropna()

def prepare_all_pairs_data(datasets, pairs):
    combined_train = pd.DataFrame()

    for pair in pairs:
        try:
            pair_data = prepare_data(datasets, pair)

            # Split data into training sets
            train_data = pair_data[(pair_data.index.year < 2018) & (pair_

            # Add Pair information
            train_data['Pair'] = f"{pair[0]}_{pair[1]}"

            combined_train = pd.concat([combined_train, train_data], ignore
        except:
            print("Failed for pair: ", pair)

    return combined_train

# Prepare features and target for train and test separately
def prepare_features_target(data):
    X = data.drop(['Target', 'Pair'], axis=1)
    y = data['Target']
    return X, y

# Handle infinities and large values
def handle_inf_nan(X):
    X = X.replace([np.inf, -np.inf], np.nan)
    X = X.fillna(method='ffill').fillna(method='bfill')
    return X

```

```

In [ ]: def train_xgboost(X_train, y_train):
    ...

    param_grid = {
        'n_estimators': [100, 300],
        'learning_rate': [0.05, 0.1],
        'max_depth': [3, 5],

```

```

'min_child_weight': [1, 3],
'gamma': [0, 0.1],
'subsample': [0.9],
'lambda': [1, 2] # L2 regularization
}

model = XGBRegressor(random_state=42)
random_search = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_grid,
    scoring='neg_mean_squared_error',
    cv=3,
    verbose=1,
    random_state=42
)

random_search.fit(X_train, y_train)

print("Best Hyperparameters:", random_search.best_params_)
#Best Hyperparameters: {'subsample': 0.9, 'n_estimators': 100, 'min_c
...
model = XGBRegressor(
    n_estimators=100,
    learning_rate=0.05,
    max_depth=3,
    min_child_weight=1,
    subsample=0.9,
    colsample_bytree=0.8,
    random_state=42,
    gamma = 0,
)
model.fit(X_train, y_train)
return model

def evaluate_model(model, X_test, y_test):
    predictions = model.predict(X_test)
    mse = mean_squared_error(y_test, predictions)
    r2 = r2_score(y_test, predictions)
    return mse, r2, predictions

def implement_trading_strategy(predictions, test_features, tolerance=0.00
    positions = np.zeros(len(predictions))
    positions[predictions > tolerance] = 1
    positions[predictions < -tolerance] = -1

    returns = positions * test_features['Spread Returns'].values
    cumulative_returns = np.cumprod(1 + returns) - 1
    return cumulative_returns, positions

```

```

In [ ]: combined_train = prepare_all_pairs_data(datasets, brim_stock_pairs)

X_train, y_train = prepare_features_target(combined_train)

```

```

# Handle infinities and large values
X_train = handle_inf_nan(X_train)

def clean_target_data(X, y):
    mask = y.notna() & ~np.isinf(y) # Create mask to filter out NaN and
    return X[mask], y[mask]

X_train, y_train = clean_target_data(X_train, y_train)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)

# Train the model
model = train_xgboost(X_train_scaled, y_train)

cumulative_returns_dict = {}
results = []

plt.figure(figsize=(12, 8))
for pair in brim_stock_pairs:
    try:
        pair_data = prepare_data(datasets, pair)

        test_data = pair_data[pair_data.index.year == 2018]
        X_test, y_test = prepare_features_target(test_data)

        # Handle infinities and large values
        X_test = handle_inf_nan(X_test)
        X_test_scaled = scaler.transform(X_test)

        mse, r2, predictions = evaluate_model(model, X_test_scaled, y_test)

        print(f"Pair: {pair}, Mean Squared Error: {mse}, R-squared: {r2}")

        # Implement trading strategy
        cumulative_returns, positions = implement_trading_strategy(predictions, pair)

        cumulative_returns_dict[pair[0]+pair[1]] = cumulative_returns
        results.append([pair[0], pair[1], mse, r2, cumulative_returns[-1]])

        plt.plot(test_data.index, cumulative_returns, label=f"{pair[0]}_{pair[1]}")

    except:
        print("failed for pair: ", pair)


# Finalize and show the plot
plt.title("Cumulative Returns for All Pairs' Trading Strategy")
plt.xlabel("Date")
plt.ylabel("Cumulative Returns")
plt.legend(loc="upper left")
plt.show()

```

Failed for pair: ['MA', 'VFC']

Pair: ['BEN', 'COG'], Mean Squared Error: 0.00451371349363909, R-squared: -0.00026810014327161014
Pair: ['DISCA', 'RIG'], Mean Squared Error: 0.001973425179299559, R-squared: -0.0031005352485402593
Pair: ['DISCK', 'RIG'], Mean Squared Error: 0.0022371665971717583, R-squared: -0.0032409391077670158
Pair: ['ADBE', 'CRM'], Mean Squared Error: 0.0009721693843248306, R-squared: -0.0027555952005899886
Pair: ['CF', 'HBI'], Mean Squared Error: 0.0016747755748894892, R-squared: -0.0029470894992782437
Pair: ['ESV', 'GNW'], Mean Squared Error: 0.001516347826730432, R-squared: -0.0006441598065316523
Pair: ['CNX', 'HBI'], Mean Squared Error: 0.12420939573188185, R-squared: -0.35479007476139124
Pair: ['AMZN', 'CRM'], Mean Squared Error: 0.0004666383504514382, R-squared: -0.005883572034045992
failed for pair: ['MA', 'VFC']
Pair: ['FCX', 'GNW'], Mean Squared Error: 0.001411615078856414, R-squared: -0.007590062015828236
Pair: ['CRM', 'NVDA'], Mean Squared Error: 0.0004768313584557387, R-squared: -0.007661109832311386
Pair: ['CF', 'FOSL'], Mean Squared Error: 0.0029678608300975212, R-squared: -0.0001652199540649324
Pair: ['FCX', 'HBI'], Mean Squared Error: 1.1685075135682383, R-squared: 0.4602091136913551
Pair: ['DISCK', 'ESV'], Mean Squared Error: 0.002576803341838784, R-squared: -0.0010491975351158978
Pair: ['DISCK', 'NE'], Mean Squared Error: 0.0007791434745997887, R-squared: -0.0032454049497929738
Pair: ['DISCA', 'NE'], Mean Squared Error: 0.0007845826742430618, R-squared: -0.003544479508166054
Pair: ['DISCA', 'ESV'], Mean Squared Error: 0.0027463758954903435, R-squared: -0.0024229656030931856
Pair: ['ESV', 'RRC'], Mean Squared Error: 0.0017591181907027987, R-squared: -1.898209434081366e-07
Pair: ['NBL', 'RIG'], Mean Squared Error: 0.0009181409687401804, R-squared: -0.0010505409766095042
Pair: ['CNX', 'GNW'], Mean Squared Error: 0.0011388236692157902, R-squared: -0.001245347539460706
Pair: ['COG', 'DO'], Mean Squared Error: 0.012672665141832372, R-squared: -0.0004250098941913638
Pair: ['HBI', 'NBL'], Mean Squared Error: 0.005846642261410964, R-squared: -0.0014558739615091465
Pair: ['HBI', 'MR0'], Mean Squared Error: 1.2656450589321255, R-squared: -0.3990049358816612
Pair: ['GNW', 'NBL'], Mean Squared Error: 0.0008294961365566999, R-squared: -0.0016159289064052729
Pair: ['DISCA', 'MA'], Mean Squared Error: 0.0004116082904857544, R-squared: -0.005828939077023199
Pair: ['DISCK', 'MA'], Mean Squared Error: 0.00040039160447767825, R-squared: -0.006019416706309455
Pair: ['RIG', 'RRC'], Mean Squared Error: 0.018031267071124398, R-squared: 0.05513578849055534

Pair: ['CF', 'CNX'], Mean Squared Error: 0.0011050683498695993, R-squared: -0.0008084589430712441
 Pair: ['CF', 'GNW'], Mean Squared Error: 0.0006406255833040104, R-squared: -0.0002880193141667764
 Pair: ['ESV', 'HBI'], Mean Squared Error: 0.002056426710537815, R-squared: 1.4162895762059868e-05
 Pair: ['NE', 'RRC'], Mean Squared Error: 0.0014078256262157335, R-squared: -0.001270143779948718
 Pair: ['ADBE', 'RHT'], Mean Squared Error: 0.005653497953014792, R-squared: -0.04248712632337526
 Pair: ['MA', 'RIG'], Mean Squared Error: 0.0003392074676467846, R-squared: -0.007404052009366335
 Pair: ['NBL', 'SWN'], Mean Squared Error: 0.0007551887186818223, R-squared: -0.0013313752361685527
 failed for pair: ['CTWS', 'AWR']
 Pair: ['CTWS', 'WTR'], Mean Squared Error: 0.001485423671635775, R-squared: -0.008006100181691123
 Pair: ['AWR', 'WTR'], Mean Squared Error: 0.0007874220880120775, R-squared: -0.013233981845165443
 Pair: ['SLB', 'PFE'], Mean Squared Error: 0.5267609816691273, R-squared: -0.5583297274057162



```
In [ ]: results_df = pd.DataFrame(results, columns=["Stock 1", "Stock 2", "MSE",
print(results_df)
```


	Stock 1	Stock 2	MSE	R-squared	Final Return
0	BEN	COG	0.004514	-2.681001e-04	-69.602278
1	DISCA	RIG	0.001973	-3.100535e-03	-26.438180
2	DISCK	RIG	0.002237	-3.240939e-03	-14.968153
3	ADBE	CRM	0.000972	-2.755595e-03	0.000000
4	CF	HBI	0.001675	-2.947089e-03	-21.584651
5	ESV	GNW	0.001516	-6.441598e-04	-17.322765
6	CNX	HBI	0.124209	-3.547901e-01	-99.999947
7	AMZN	CRM	0.000467	-5.883572e-03	0.000000
8	FCX	GNW	0.001412	-7.590062e-03	-22.036002
9	CRM	NVDA	0.000477	-7.661110e-03	0.000000
10	CF	FOSL	0.002968	-1.652200e-04	-62.163850
11	FCX	HBI	1.168508	4.602091e-01	-99.687298
12	DISCK	ESV	0.002577	-1.049198e-03	-67.246877
13	DISCK	NE	0.000779	-3.245405e-03	0.000000
14	DISCA	NE	0.000785	-3.544480e-03	0.000000
15	DISCA	ESV	0.002746	-2.422966e-03	-69.890783
16	ESV	RRC	0.001759	-1.898209e-07	-39.783494
17	NBL	RIG	0.000918	-1.050541e-03	0.000000
18	CNX	GNW	0.001139	-1.245348e-03	0.000000
19	COG	DO	0.012673	-4.250099e-04	-99.237835
20	HBI	NBL	0.005847	-1.455874e-03	-92.251316
21	HBI	MRO	1.265645	-3.990049e-01	-99.993541
22	GNW	NBL	0.000829	-1.615929e-03	-5.929722
23	DISCA	MA	0.000412	-5.828939e-03	0.000000
24	DISCK	MA	0.000400	-6.019417e-03	0.000000
25	RIG	RRC	0.018031	5.513579e-02	-100.036212
26	CF	CNX	0.001105	-8.084589e-04	-11.124088
27	CF	GNW	0.000641	-2.880193e-04	0.000000
28	ESV	HBI	0.002056	1.416290e-05	-38.919603
29	NE	RRC	0.001408	-1.270144e-03	0.000000
30	ADBE	RHT	0.005653	-4.248713e-02	-91.495293
31	MA	RIG	0.000339	-7.404052e-03	0.000000
32	NBL	SWN	0.000755	-1.331375e-03	0.000000
33	CTWS	WTR	0.001485	-8.006100e-03	0.000000
34	AWR	WTR	0.000787	-1.323398e-02	0.000000
35	SLB	PFE	0.526761	-5.583297e-01	-100.664589

```
In [ ]: print(results_df[results_df['Final Return']>0].reset_index())
```


Empty DataFrame

Columns: [index, Stock 1, Stock 2, MSE, R-squared, Final Return]

Index: []

```
In [ ]: # Feature importance
feature_importance = model.feature_importances_
sorted_idx = np.argsort(feature_importance)
pos = np.arange(sorted_idx.shape[0]) + .5

plt.figure(figsize=(12, 6))
plt.barh(pos, feature_importance[sorted_idx], align='center')
plt.yticks(pos, X_train.columns[sorted_idx])
plt.title('Feature Importance')
plt.show()
```



Other models

```
In [ ]: def prepare_all_pairs_data(datasets, pairs, start_year, end_year):
    combined_data = pd.DataFrame()

    for pair in pairs:
        try:
            pair_data = prepare_data(datasets, pair)

            # Filter data for the specified year range
            pair_data = pair_data[(pair_data.index.year >= start_year) &
                                   (pair_data.index.year <= end_year)]

            # Add Pair information
            pair_data['Pair'] = f"{pair[0]}_{pair[1]}"

            combined_data = pd.concat([combined_data, pair_data], ignore_index=True)
        except Exception as e:
            print(f"Failed for pair: {pair}. Error: {e}")

    return combined_data

def train_and_evaluate_model(model, X_train, y_train, X_test, y_test, test_size=0.2):
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    mse = mean_squared_error(y_test, predictions)
    r2 = r2_score(y_test, predictions)
    cumulative_returns, positions = implement_trading_strategy(predictions, y_test)
    return mse, r2, cumulative_returns, positions
```

```
In [ ]: # Prepare combined training data (2014-2017)
combined_train = prepare_all_pairs_data(datasets, brim_stock_pairs, 2014, 2017)
X_train, y_train = prepare_features_target(combined_train)

# Handle infinities and large values
```

```

X_train = handle_inf_nan(X_train)
X_train, y_train = clean_target_data(X_train, y_train)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)

```

Failed for pair: ['MA', 'VFC']. Error: 'VFC'

```

In [ ]: models = {
    'XGBoost': xgb.XGBRegressor(n_estimators=100, learning_rate=0.05, max
                                min_child_weight=1, subsample=0.9, colsam
                                random_state=42, gamma=0),
    'Linear Regression': LinearRegression()
}
# 'Linear Regression': LinearRegression()
trained_models = {name: model.fit(X_train_scaled, y_train) for name, mode

```

```

In [ ]: results = {model_name: {'results': [], 'cumulative_returns': {}} for mode

for pair in brim_stock_pairs:
    try:
        pair_data = prepare_data(datasets, pair)
        test_data = pair_data[pair_data.index.year == 2018]
        X_test, y_test = prepare_features_target(test_data)

        # Handle infinities and large values
        X_test = handle_inf_nan(X_test)
        X_test_scaled = scaler.transform(X_test)

        for model_name, model in trained_models.items():
            mse, r2, cumulative_returns, positions = train_and_evaluate_m
                model, X_train_scaled, y_train, X_test_scaled, y_test, te
            )

            results[model_name]['results'].append([pair[0], pair[1], mse,
            results[model_name]['cumulative_returns'][f"{pair[0]}_{pair[1]

        except Exception as e:
            print(f"Failed for pair: {pair}. Error: {e}")

```

Failed for pair: ['MA', 'VFC']. Error: 'VFC'

Failed for pair: ['CTWS', 'AWR']. Error: Input contains infinity or a value too large for dtype('float64').

```


In [ ]: # Convert results to DataFrames
for model_name in results:
    results[model_name]['results_df'] = pd.DataFrame(
        results[model_name]['results'],
        columns=["Stock 1", "Stock 2", "MSE", "R-squared", "Final Return"]
    )
for model_name, model_results in results.items():
    cumulative_returns = model_results['cumulative_returns']


```

```
plt.figure(figsize=(15, 10))

for pair, returns in cumulative_returns.items():
    plt.plot(test_data.index, returns, label=f"{pair}")

plt.title(f"Cumulative Returns for All Pairs' Trading Strategy ({mode})")
plt.xlabel("Date")
plt.ylabel("Cumulative Returns")
plt.legend(loc="upper left", bbox_to_anchor=(1, 1))
plt.tight_layout()
plt.show()
```





```
In [ ]: # Print results for all models
for model_name, model_results in results.items():
    print(f"\nResults for {model_name}:")
    print(model_results['results_df'])
    print(f"\nPositive returns for {model_name}:")
    print(model_results['results_df'][model_results['results_df']['Final
```

Results for XGBoost:

	Stock 1	Stock 2	MSE	R-squared	Final Return
0	BEN	COG	0.004514	-2.681001e-04	-69.602278
1	DISCA	RIG	0.001973	-3.100535e-03	-26.438180
2	DISCK	RIG	0.002237	-3.240939e-03	-14.968153
3	ADBE	CRM	0.000972	-2.755595e-03	0.000000
4	CF	HBI	0.001675	-2.947089e-03	-21.584651
5	ESV	GNW	0.001516	-6.441598e-04	-17.322765
6	CNX	HBI	0.124209	-3.547901e-01	-99.999947
7	AMZN	CRM	0.000467	-5.883572e-03	0.000000
8	FCX	GNW	0.001412	-7.590062e-03	-22.036002
9	CRM	NVDA	0.000477	-7.661110e-03	0.000000
10	CF	FOSL	0.002968	-1.652200e-04	-62.163850
11	FCX	HBI	1.168508	4.602091e-01	-99.687298
12	DISCK	ESV	0.002577	-1.049198e-03	-67.246877
13	DISCK	NE	0.000779	-3.245405e-03	0.000000
14	DISCA	NE	0.000785	-3.544480e-03	0.000000
15	DISCA	ESV	0.002746	-2.422966e-03	-69.890783
16	ESV	RRC	0.001759	-1.898209e-07	-39.783494
17	NBL	RIG	0.000918	-1.050541e-03	0.000000
18	CNX	GNW	0.001139	-1.245348e-03	0.000000
19	COG	DO	0.012673	-4.250099e-04	-99.237835
20	HBI	NBL	0.005847	-1.455874e-03	-92.251316

21	HBI	MRO	1.265645	-3.990049e-01	-99.993541
22	GNW	NBL	0.000829	-1.615929e-03	-5.929722
23	DISCA	MA	0.000412	-5.828939e-03	0.000000
24	DISCK	MA	0.000400	-6.019417e-03	0.000000
25	RIG	RRC	0.018031	5.513579e-02	-100.036212
26	CF	CNX	0.001105	-8.084589e-04	-11.124088
27	CF	GNW	0.000641	-2.880193e-04	0.000000
28	ESV	HBI	0.002056	1.416290e-05	-38.919603
29	NE	RRC	0.001408	-1.270144e-03	0.000000
30	ADBE	RHT	0.005653	-4.248713e-02	-91.495293
31	MA	RIG	0.000339	-7.404052e-03	0.000000
32	NBL	SWN	0.000755	-1.331375e-03	0.000000
33	CTWS	WTR	0.001485	-8.006100e-03	0.000000
34	AWR	WTR	0.000787	-1.323398e-02	0.000000
35	SLB	PFE	0.526761	-5.583297e-01	-100.664589

Positive returns for XGBoost:

Empty DataFrame

Columns: [Stock 1, Stock 2, MSE, R-squared, Final Return]

Index: []

Results for Linear Regression:

	Stock 1	Stock 2	MSE	R-squared	Final Return
0	BEN	COG	0.004511	0.000289	2.436705e+01
1	DISCA	RIG	0.001988	-0.010529	-6.000638e+01
2	DISCK	RIG	0.002252	-0.009949	-6.304569e+01
3	ADBE	CRM	0.000987	-0.017856	-1.749857e+00
4	CF	HBI	0.001684	-0.008622	-4.296583e+01
5	ESV	GNW	0.001513	0.001642	1.981945e+02
6	CNX	HBI	0.092709	-0.011203	4.922263e+06
7	AMZN	CRM	0.000480	-0.034224	-3.816226e+01
8	FCX	GNW	0.001403	-0.001276	9.610367e+01
9	CRM	NVDA	0.000486	-0.027352	8.802548e+00
10	CF	FOSL	0.002975	-0.002636	5.935695e+01
11	FCX	HBI	2.177121	-0.005719	2.044167e+10
12	DISCK	ESV	0.002582	-0.003177	-6.743625e+00
13	DISCK	NE	0.000786	-0.012492	-3.308104e+01
14	DISCA	NE	0.000792	-0.012781	-3.364055e+01
15	DISCA	ESV	0.002748	-0.002949	-1.103411e+01
16	ESV	RRC	0.001757	0.001293	1.825549e+02
17	NBL	RIG	0.000916	0.001064	4.829557e+01
18	CNX	GNW	0.001138	-0.000181	2.786207e+01
19	COG	DO	0.012780	-0.008902	2.704250e+02
20	HBI	NBL	0.005851	-0.002220	-1.123061e+01
21	HBI	MRO	0.904465	0.000233	1.221062e+20
22	GNW	NBL	0.000830	-0.001839	4.973311e+01
23	DISCA	MA	0.000425	-0.038245	-2.792092e+01
24	DISCK	MA	0.000414	-0.039504	-2.788354e+01
25	RIG	RRC	0.019146	-0.003279	1.215596e+03
26	CF	CNX	0.001113	-0.007997	-2.717591e+01
27	CF	GNW	0.000644	-0.005875	-2.602832e+00
28	ESV	HBI	0.002053	0.001591	2.390695e+02
29	NE	RRC	0.001407	-0.001035	2.633550e+01

30	ADBE	RHT	0.005420	0.000524	2.427129e+02
31	MA	RIG	0.000347	-0.031217	7.110779e+00
32	NBL	SWN	0.000754	0.000522	2.682372e+01
33	CTWS	WTR	0.001502	-0.019102	-3.317890e+01
34	AWR	WTR	0.000804	-0.034022	-5.343680e+01
35	SLB	PFE	0.337668	0.001067	1.627211e+07

Positive returns for Linear Regression:

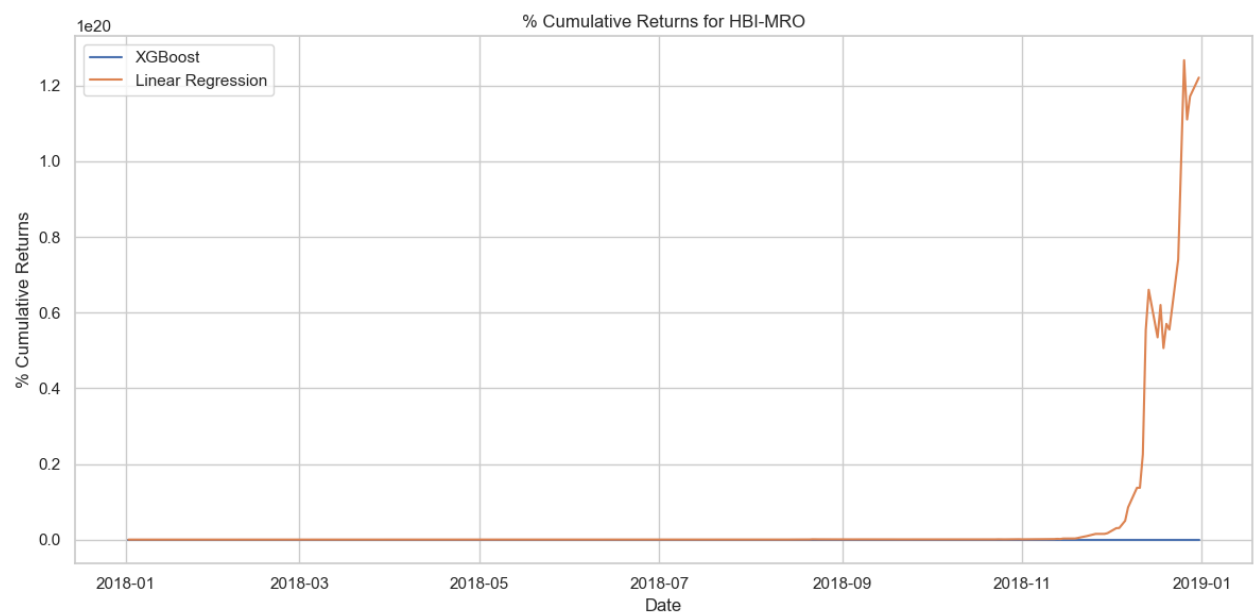


	Stock 1	Stock 2	MSE	R-squared	Final Return
0	BEN	COG	0.004511	0.000289	2.436705e+01
1	ESV	GNW	0.001513	0.001642	1.981945e+02
2	CNX	HBI	0.092709	-0.011203	4.922263e+06
3	FCX	GNW	0.001403	-0.001276	9.610367e+01
4	CRM	NVDA	0.000486	-0.027352	8.802548e+00
5	CF	FOSL	0.002975	-0.002636	5.935695e+01
6	FCX	HBI	2.177121	-0.005719	2.044167e+10
7	ESV	RRC	0.001757	0.001293	1.825549e+02
8	NBL	RIG	0.000916	0.001064	4.829557e+01
9	CNX	GNW	0.001138	-0.000181	2.786207e+01
10	COG	DO	0.012780	-0.008902	2.704250e+02
11	HBI	MRO	0.904465	0.000233	1.221062e+20
12	GNW	NBL	0.000830	-0.001839	4.973311e+01
13	RIG	RRC	0.019146	-0.003279	1.215596e+03
14	ESV	HBI	0.002053	0.001591	2.390695e+02
15	NE	RRC	0.001407	-0.001035	2.633550e+01
16	ADBE	RHT	0.005420	0.000524	2.427129e+02
17	MA	RIG	0.000347	-0.031217	7.110779e+00
18	NBL	SWN	0.000754	0.000522	2.682372e+01
19	SLB	PFE	0.337668	0.001067	1.627211e+07



```
In [ ]: # Separate plots for pairs in the paper
specific_pairs = [
    ('CNX', 'HBI'), ('FCX', 'HBI'), ('HBI', 'MRO'),
    ('ESV', 'RRC'), ('ESV', 'GNW')
]

for pair in specific_pairs:
    plt.figure(figsize=(12, 6))
    pair_key = f"{pair[0]}_{pair[1]}"


    for model_name, model_results in results.items():
        if pair_key in model_results['cumulative_returns']:
            returns = model_results['cumulative_returns'][pair_key]
            plt.plot(test_data.index, returns, label=model_name)

    plt.title(f"% Cumulative Returns for {pair[0]}-{pair[1]}")
    plt.xlabel("Date")
    plt.ylabel("% Cumulative Returns")
    plt.legend()
    plt.tight_layout()
    plt.show()
```

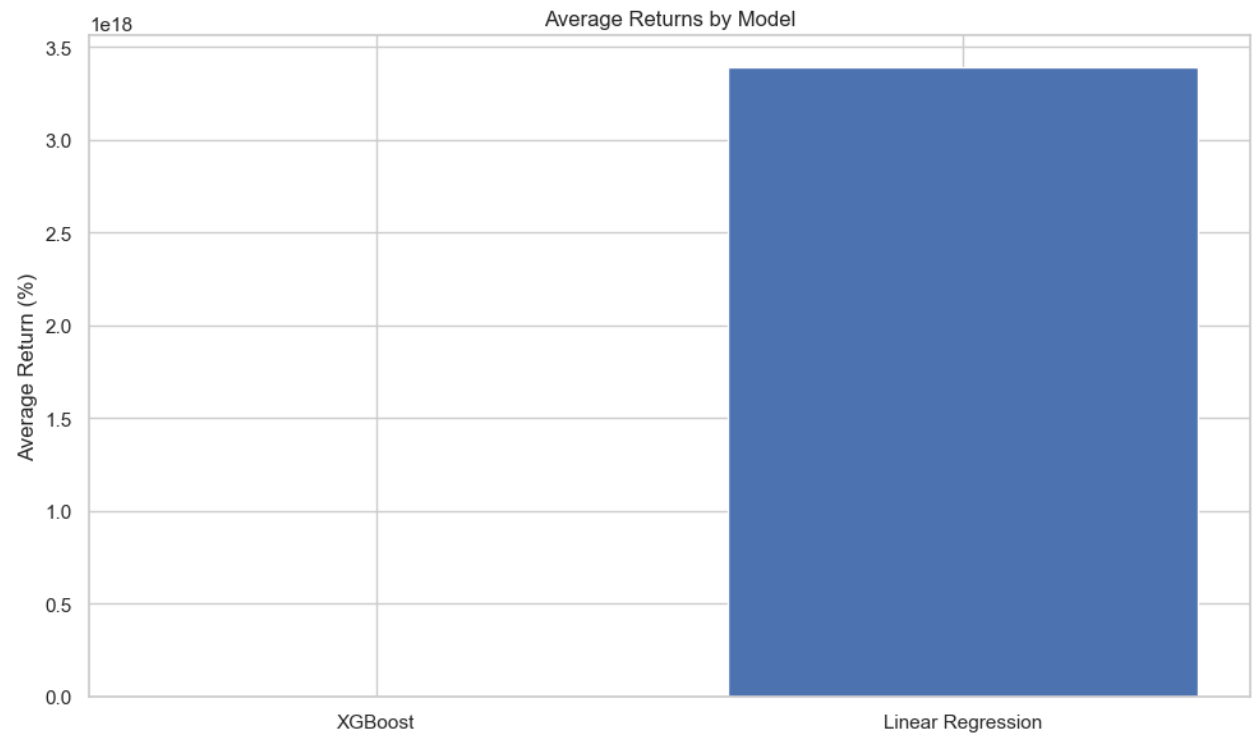




```
In [ ]: # Heatmap of correlation between features
plt.figure(figsize=(12, 10))
sns.heatmap(X_train.corr(), annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap of Features')
plt.tight_layout()
plt.show()
```



```
In [ ]: # Bar plot of average returns by model
avg_returns = {model: results[model]['results_df']['Final Return'].mean()}
plt.figure(figsize=(10, 6))
plt.bar(avg_returns.keys(), avg_returns.values())
plt.title('Average Returns by Model')
plt.ylabel('Average Return (%)')
plt.tight_layout()
plt.show()
```



```
In [ ]: # Table of top 5 performing pairs for each model
for model_name, model_results in results.items():
    top_5 = model_results['results_df'].nlargest(5, 'Final Return')
    print(f"\nTop 5 performing pairs for {model_name}:")
    print(tabulate(top_5, headers='keys', tablefmt='pretty', floatfmt=".2"))
```

Top 5 performing pairs for XGBoost:

	Stock 1	Stock 2	MSE	R-squared
Final Return				
3	ADBE	CRM	0.0009721693843248306	-0.0027555952005899886
0.0				
7	AMZN	CRM	0.0004666383504514382	-0.005883572034045992
0.0				
9	CRM	NVDA	0.0004768313584557387	-0.007661109832311386
0.0				
13	DISCK	NE	0.0007791434745997887	-0.0032454049497929738
0.0				
14	DISCA	NE	0.0007845826742430618	-0.003544479508166054
0.0				

Top 5 performing pairs for Linear Regression:

	Stock 1	Stock 2	MSE	R-squared
Final Return				
21	HBI	MRO	0.9044646688384439	0.00023270576098088913
1.2210619658254595e+20				
11	FCX	HBI	2.177120673630198	-0.0057187346883840195
20441670124.75649				
35	SLB	PFE	0.3376684340646472	0.0010673206239936173
16272107.207067419				
6	CNX	HBI	0.09270880925149369	-0.011203491304254465
4922263.355409072				
25	RIG	RRC	0.01914602861346223	-0.003279312432044712
1215.5955451655107				